


Week1 &2 lately report

Qian Jiang

1. Finished reading all papers about Xnor-Net & Summary of quantized network
2. Finished studying Neural Network related knowledge including perceptron, backpropagation, CNN, BNN...
(Notes are attched below)
3. Understood codes of XNOR-Net/util.py
4. Run codes (met lots of bugs, thanks to Weitian's help)

Neural Network Basics

1. M-P Model




2. Perceptron

Diagram: Input x_i is multiplied by weight w_i and summed (Σ) with bias b to produce a weighted sum. This sum is then passed through an activation function f to produce the output y . The equation is $y = f(\sum w_i x_i + b)$ (for: linear separation).

↓

Multi-layer Perceptron.

3. Back Propagation



$$a_j^L = \sigma(\sum_k w_{jk}^L a_k^{L-1} + b_j^L)$$

$$a^L = \sigma(\sum_k w_{jk}^L a_k^{L-1} + b_j^L)$$

z^L : weighted input

Back-info:

$$1. S \odot t = S_i t_i$$

$$\text{eg } \begin{pmatrix} 1 \\ 2 \end{pmatrix} \odot \begin{pmatrix} 3 \\ 4 \end{pmatrix}$$

$$= \begin{pmatrix} 3 \\ 8 \end{pmatrix}$$

Hadamard Product.

Cost function (quadratic): $C = \frac{1}{2n} \sum_x |y(x) - a^L(x)|^2$

$$C_x = \frac{1}{2} |y - a|^2$$

Equations of BP

$$1. \delta^L = \nabla_a C \odot \sigma'(z^L)$$

(L: Last)

$$\left(\delta_j^L = \frac{\partial C}{\partial z_j^L} = \frac{\partial C}{\partial a_j^L} \cdot \sigma'(z_j^L) \right)$$

$$\left(\delta^L = \nabla_a C \odot \sigma'(z^L) \right)$$

$z^{L+1} = f(z^L)$ chain rule

$$2. \nabla_a C:$$

matrix for

$$\frac{\partial C}{\partial a_j^L}$$

$$2. \delta^L = (W^{L+1})^T \cdot \delta^{L+1} \odot \sigma'(z^L)$$

$$3. \frac{\partial C}{\partial b_j^L} = \delta_j^L$$

$$4. \frac{\partial C}{\partial w_{jk}^L} = a_k^{L-1} \delta_j^L$$

$$\delta_j^L = \frac{\partial C}{\partial z_j^L} = \sum_k \frac{\partial C}{\partial z_k^{L+1}} \cdot \frac{\partial z_k^{L+1}}{\partial z_j^L}$$

$$z_k^{L+1} = \sum_i w_{ki}^{L+1} a_i^L + b_k^{L+1} = \sum_{ki} w_{ki}^{L+1} \sigma(z_i^L) + b_k^{L+1}$$

$$\frac{\partial z_k^{L+1}}{\partial z_j^L} = w_{kj}^{L+1} \sigma'(z_j^L)$$

$$(z^L = w^L a^{L-1} + b^L)$$

gradient
vanishing:

for ③ $\frac{\partial C}{\partial b_j^L} = \delta_j^L$

$$\frac{\partial C}{\partial b_j^L} = \frac{\partial C}{\partial z_j^L} \cdot \underbrace{\frac{\partial z_j^L}{\partial b_j^L}}_1 = \frac{\partial C}{\partial z_j^L} = \delta_j^L$$


e.g. sigmoid

for ④ $\frac{\partial C}{\partial w_{jk}^L} = a_k^{L-1} \delta_j^L$ $\frac{\partial C}{\partial w_{jk}^L} = \frac{\partial C}{\partial z_j^L} \cdot \frac{\partial z_j^L}{\partial w_{jk}^L}$

Also $\frac{\partial C}{\partial w} = a_{in} \delta_{out}$

$$= \delta_j^L \cdot a_k^{L-1}$$

↓ derivation



★ when $a_{in} \rightarrow 0$ or output is saturated (high or low)
 $\delta_{out} \rightarrow 0$



$$\frac{\partial C}{\partial b_i} = \frac{\partial C}{\partial a_L} \frac{\partial a_L}{\partial b_i}$$

Matrix Form:

$$\delta^L = \Sigma'(z^L) \nabla a_C$$

$$\Sigma'(z^L) : \begin{pmatrix} \sigma'(z_1^L) & 0 & \dots \\ 0 & \sigma'(z_2^L) & \dots \\ \vdots & \vdots & \ddots \end{pmatrix}$$

$$\delta^L = \Sigma'(z^L) (W^{L+1})^T \delta^{L+1}$$

↓ smaller
→ vanish

gradient

exploding:

$$\sigma' w > 1$$

BP Algorithm:

1. Input x
2. Feedforward : $z^L = w^L a^{L-1} + b^L$, $a^L = \sigma(z^L)$
3. Output error : $\delta^L = \nabla a_C \cdot \sigma'(z^L)$
4. Backpropagate the error : $\delta^L = ((W^{L+1})^T \delta^{L+1}) \odot \sigma'(z^L)$
5. Output : $\frac{\partial C}{\partial b_j^L} = \delta_j^L$ $\frac{\partial C}{\partial w_{jk}^L} = a_k^{L-1} \delta_j^L$

Gradient descent:

update : $w^L \rightarrow w^L - \frac{\eta}{m} \delta^L (a^{L-1})^T$

$b^L \rightarrow b^L - \frac{\eta}{m} \delta^L$

m : batch size

η : learning rate


$$\Delta C = \frac{\partial C}{\partial a_m^L} \cdot \frac{\partial a_m^L}{\partial a_n^{L-1}} \cdot \frac{\partial a_n^{L-1}}{\partial a_p^{L-2}} \dots \frac{\partial a_j^L}{\partial w_{jk}^L} \cdot \Delta w_{jk}^L$$

↓
 Σ all possible path

Convolutional Neural Networks. (stanford CS231n)

CONV \rightarrow Relu \rightarrow pool - Fully-Connected

eg. Input $[32 \times 32 \times 3]$ \Rightarrow CONV $[32 \times 32 \times 12]$ \Rightarrow RELU $[32 \times 32 \times 12]$
 image size \downarrow 3 color channel \downarrow 12 filter \downarrow activation function applied
 \downarrow
 FC $[1 \times 1 \times 10]$ \Leftarrow POOL $[16 \times 16 \times 12]$
 \downarrow 10 class scores \downarrow downsampling



Parameter

1. CONV

Sharing:

Filter

remain same as input slide

F.F.D. weight per filter

depth: number of filters

stride: move the filter S pixel at a time

zero-padding: zeros around the input border

W : input volume size

F : filter size

P : amount of zero padding

S : stride

K : Number of filters

input $W_1 \times H_1 \times D_1$

output $W_2 = (W_1 - F + 2P) / S + 1$

$H_2 = (H_1 - F + 2P) / S + 1$

$D_2 = K$

eg. $4 \times 4 \times 1$ (pad 1)

0	0	0	0
0	1	2	0
0	3	4	0
0	0	0	0

Filter $2 \times 2 \times 1$

$S = 2$

$*$ $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$3 \ 4$

$b_0 = 0$

$b_0 + 0 \times 1 + 3 \times 2 + 0 \times 3 + 0 \times 4 = 6$

$\begin{bmatrix} 4 & 6 \\ 6 & 4 \end{bmatrix}$

2. ReLU



3. Pooling: $W_1 \times H_1 \times D_1$

F : spatial extent

S : stride.

$$W_2 = (W_1 - F) / S + 1$$

$$H_2 = (H_1 - F) / S + 1$$

$$D_2 = D_1$$

e.g.

1 1 2 4

5 6 7 8

3 2 1 0

1 2 3 4

max pool

2×2 filter

stride 2.

6 8

3 4

Back propagation

4. (Normalization layer: optional)

5. Fully - Connected Layer:

a layer have full connections to all activations in the previous layer.

(similar to normal Neural Networks)

Some common CNN:

LeNet, AlexNet, ZFNet, GoLeNet, VGGNet, ResNet.

Paper Notes : CVPR 16 XNOR-Net: ImageNet Class.

top-1 score: XNOR-Net : ImageNet Classification ... (CVPR 16)


the top class Ab: BNN : weight binarized 32x memory saving
(highest proba) XNOR-Net : both input & weight binarized, FC binarized
is the same 58x faster 32x memory saving

as the target run networks on CPUs
label classification accuracy : BWN of Alexnet is same as full-precision Alexnet

top-5 score: compare with BinaryConnect, Binary Nets.
the target 1b% accuracy (top-1)

1. Intro	operation	Mem-saving	Computation saving	Accu (Alexnet)
standard	+ + . - . X	1x	1x	56.7
BW	+ + . -	~32x	~2x	56.8
BW-BI	XNOR.bitcount	~32x	~58x	44.2

XNOR :



a b f

0 0 1

0 1 0

1 0 0

1 1 1

2. Related Work

1. shallow networks:

estimating a deep NN with a shallower model

pros: network size ↓

cons: difficult to train with large number of parameters

2. Compressing pre-trained deep networks:

pruning redundant non-informative weights

remove redundant weights. Huffman coding

HashNets. Matrix factorization

loss function

= cost function

3. Quantizing parameters:

quantize the weights of FC layer by vector quantization

Hash function:

map data

to fixed size

4. Network binarization:

binarize weights & activations

3. Binary CNN ① BWN

I_L : input tensor for the L th layer

W_{lk} : the k -th weight filter in the L th layer.

K^L : the number of weight filter in the L th layer.

$I \in \mathbb{R}^{c \times w_{in} \times h_{in}}$ (c : channel)

$W \in \mathbb{R}^{c \times w \times h}$

Binary filter $B \in \{+1, -1\}^{c \times w \times h}$

scaling factor α : $W \approx \alpha B$

$$I * W = (I \oplus B) \alpha$$

(1)

↳ convolution without multiplication

$$W_{lk} = A_{lk} B_{lk}$$

scalar

binary tensor

optimize $W \approx \alpha B$:

$$J(B, \alpha) = \|W - \alpha B\|^2 \quad (2)$$

$$= \alpha^2 \underbrace{B^T B}_{\text{constant } n} - 2\alpha \underbrace{W^T B}_{\text{constant } C} + \underbrace{W^T W}_{\text{constant } C} \quad (3)$$

$$= \alpha^2 n - 2\alpha W^T B + C$$

$$B^* = \arg \max_B \{W^T B\} \Rightarrow B_i = \begin{cases} +1 & \text{if } W_i \geq 0 \\ -1 & \text{if } W_i < 0 \end{cases} \quad (4)$$

$$\therefore B^* = \text{sign}(W)$$

$$\frac{\partial J}{\partial \alpha} = 2\alpha n - 2W^T B = 0 \Rightarrow \alpha^* = \frac{W^T B^*}{n} \quad (5)$$

$$\therefore \alpha^* = \frac{W^T \text{sign}(W)}{n} = \frac{\sum |W_i|}{n} = \frac{1}{n} \|W\|_1 \quad (6)$$

Train BWN:

$$\text{for } \text{sign}(r) \quad \frac{\partial \text{sign}}{\partial r} = r \text{ } (|r| \leq 1)$$

$$\frac{\partial C}{\partial W_i} = \frac{\partial C}{\partial W_i} \left(\frac{1}{n} + \frac{\partial \text{sign}}{\partial W_i} \alpha \right) ?$$

key: ① only binarize the weights during forward & bp
② for update, use the high precision (real-value) weight.

② XNOR - Network

bitcount:
count the
number of
set bits in
a string

Key: 1. binarize both weights & inputs

2. convolution can be implemented by XOR & bitcount

Binary dot product:

$$X^T W \approx \beta H^T \alpha B, H, B \in \{+1, -1\}^n, \beta, \alpha \in \mathbb{R}^+$$

$$\text{optimize: } \alpha^*, \beta^*, H^*, B^* = \arg \min_{\alpha, \beta, H, B} \|X \odot W - \beta \alpha H \odot B\|$$

$$Y \in \mathbb{R}^n, Y_i = X_i W_i$$

$$C \in \{+1, -1\}^n, C_i = H_i B_i$$

$$r \in \mathbb{R}^+, r = \beta \alpha$$

$$\therefore Y^*, C^* = \arg \min_{Y, C} \|Y - r C\|$$

$$C^* = \text{sign}(Y) = \text{sign}(X) \odot \text{sign}(W) = H^* \odot B^*$$

$$\therefore E[|Y_i|] = E[|X_i| |W_i|] = E[|X_i|] E[|W_i|]$$

$$\therefore r^* = \frac{\sum |Y_i|}{n} = \frac{\sum |X_i| |W_i|}{n}$$

$$\approx \left(\frac{1}{n} \|X\|_{L_1}\right) \left(\frac{1}{n} \|W\|_{L_1}\right) = \beta^* \alpha^*$$

Binary convolution:

$$A = \frac{\sum |I_{1:n, i}|}{c}$$

$$K = A * k$$

\hookrightarrow 2D filter, $\forall k_{ij} = \frac{1}{w \times h}$

k_{ij} correspond to β for a sub-tensor centered at i, j

$$\therefore I * W \approx (\text{sign}(I) \otimes \text{sign}(W)) \odot K \alpha$$

\hookrightarrow bitcount & XNOR operation

Train XNOR-Net:

BNorm \rightarrow BinActiv \rightarrow Binconv \rightarrow pool


Gradient: binarize g_{in} in backpass.

use $\max_i(|g_{in}^i|)$ as scaling factor

Common Quantized Model (from paperweekly.blog)

A-function. 1. Deep Compression


1. Sign



Pruning: — Quantization — Huffman Encoding
less number of weights less bits per weight

2. tanh

$$y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



2. Binary-Net

• Quantized Neural Networks:


Training Neural Networks with low precision weights & activations

• Binarized Neural Networks:

Training deep Neural Networks with W & A constrained to ± 1

• XNOR-Net: ImageNet classification using binary convolutional neural networks

3. Sigmoid




BNN: Binarization:


$$\textcircled{1} x^b = \text{sign}(x) = \begin{cases} +1 & \text{if } x > 0 \\ -1 & \text{otherwise} \end{cases}$$

$$\textcircled{2} x^b = \text{sign}(x) = \begin{cases} +1 & \text{probability } p(x) \\ -1 & 1-p \end{cases}$$

Sign



4. ReLU



$$\sigma(x) = \text{clip}\left(\frac{x+1}{2}, 0, 1\right) = \max\left(0, \min\left(\frac{x+1}{2}\right)\right)$$

$$q = \text{sign}(r)$$

$$q_r = q \cdot |r| \leq 1 \quad ??$$

Algorithm: 1. forward propagation:

Batch Norm:

$$\text{for } k=1 \text{ to } L: W_k^b \leftarrow \text{Binarize}(W_k); S_k \leftarrow A_{k-1}^b W_k^b;$$

shifting inputs

$$A_k \leftarrow \text{Batch Norm}(S_k, A_k)$$

to: zero-mean

$$\text{if } k < L, A_k^b \leftarrow \text{Binarize}(A_k)$$

unit variance

$$2. \text{BP: for } k=L-1: \text{ if } k < L, g_{A_k} \leftarrow g_{A_k} \text{ or } |A_k| < 1$$

$$g_{A_{k-1}}^b \leftarrow g_{S_k} W_k^b$$

$$g_{W_k}^b \leftarrow g_{S_k}^T A_{k-1}^b$$

Time

Complexity:

eg.

$$t = 5n^3 + 2n$$

Time complexity

$$= O(n^3)$$

3. Accumulating the parameters gradient

for $k=1$ to L : $\Theta_k^{t+1} \leftarrow \text{update}(\Theta_k, \eta, g_{\Theta_k})$

$$W_k^{t+1} \leftarrow \text{Clip}(\text{Update}(W_k, \eta, g_{W_k}), -1, 1)$$

$$\eta^{t+1} \leftarrow \lambda \eta$$

(SBN)

BN \rightarrow shift-based Batch Normalization (no matrix multiplication)



shift-based AdaMax

Xnor-net : factor α

$$I * W = (I \oplus B) \alpha$$